

复旦大学大数据学院
School of Data Science, Fudan University

魏忠钰

Adversarial Search


April 4th, 2018

Types of Games

- Many different kinds of games!

- Axes:

- **Deterministic** or stochastic?
- One, **two**, or more players?
- **Zero sum**?
- **Perfect information** (can you see the state)?




- Want algorithms for calculating a **strategy (policy)** which recommends a move from each state

Deterministic Games

- Many possible formalizations, one is:
 - S : states (start at s_0)
 - Player (s): the player has the move in this state
 - Actions (s): A set of legal moves in a state
 - Results (s, a): A transition model, return the results of a move
 - Terminal Test (s): {true, false} if s is the terminal state
 - Terminal Utilities (s, p): A utility function gives the final numeric value of a game
- Solution for a player is a **policy**: $S \rightarrow A$ (a set of actions)
- What is the policy for depth-first-search?

Zero-Sum Games



■ Zero-Sum Games


- Agents have opposite utilities (values on outcomes)
- Lets us think of a single value that one maximizes and the other minimizes
- Adversarial, pure competition




■ General Games

- Agents have independent utilities (values on outcomes)
- Cooperation, indifference, competition, and more are all possible

Adversarial Search



Single-Agent Trees




Value of a State

Value of a state:
The best achievable
outcome (utility)
from that state

Non-Terminal States:


$$V(s) = \max_{s' \in \text{children}(s)} V(s')$$



Terminal States:

$$V(s) = \text{known}$$

Adversarial Game Trees




Minimax Values

States Under Agent's Control:

$$V(s) = \max_{s' \in \text{successors}(s)} V(s')$$

States Under Opponent's Control:

$$V(s') = \min_{s \in \text{successors}(s')} V(s)$$



Terminal States:

$$V(s) = \text{known}$$

Tic-Tac-Toe Game Tree



MAX (X)



MIN (O)





MAX (X)



MIN (O)


TERMINAL

Utility



Adversarial Search (Minimax)

- Deterministic, zero-sum games:
 - Tic-tac-toe, chess
 - One player maximizes result
 - The other minimizes result
- Minimax search:
 - A state-space search tree
 - Players alternate turns
 - Compute each node's **minimax value**: the best achievable utility against a rational (optimal) adversary



Minimax Implementation

```
def max-value(state):
```


```
    initialize v =  $-\infty$ 
```

```
    for each successor of state:
```

```
        v = max(v, min-value(successor))
```

```
    return v
```

$$V(s) = \max_{s' \in \text{successors}(s)} V(s')$$



```
def min-value(state):
```

```
    initialize v =  $+\infty$ 
```

```
    for each successor of state:
```

```
        v = min(v, max-value(successor))
```

```
    return v
```

$$V(s') = \min_{s \in \text{successors}(s')} V(s)$$


Minimax Implementation (Dispatch)

```
def value(state):
```

if the state is a terminal state: return the state's utility

if the next agent is **MAX**: return **max-value(state)**

if the next agent is **MIN**: return **min-value(state)**



```
def max-value(state):
```

initialize $v = -\infty$

for each successor of state:

$v = \max(v,$
 $\text{value(successor)})$

return v

```
def min-value(state):
```


initialize $v = +\infty$

for each successor of state:

$v = \min(v,$
 $\text{value(successor)})$


return v

Minimax Example



- How efficient is minimax?



- Just like (exhaustive) DFS
- Time: $O(b^m)$
- Space: $O(bm)$




- Example: For chess, $b \approx 35$, $m \approx 100$

- Exact solution is completely infeasible
- But, do we need to explore the whole tree?

Minimax Properties




Optimal against a perfect player. Otherwise?



Resource Limits


- Problem: In realistic games, cannot search to leaves!
- Solution: Depth-limited search
 - Search only to a limited depth in the tree
 - Need an evaluation function for non-terminal positions
- Guarantee of optimal play is gone
- More steps forward makes a BIG difference
- Use iterative deepening for an anytime algorithm



- Evaluation functions are always imperfect
- The deeper in the tree the evaluation function is buried, the less the quality of the evaluation function matters
- **It takes time to compute the evaluation function.**
 - An important example of the tradeoff between complexity of features and complexity of computation

Evaluation Functions

- Evaluation functions score non-terminals in depth-limited search



- Ideal function: returns the actual minimax value of the position


$$Eval(s) = w_1 f_1(s) + w_2 f_2(s) + \dots + w_n f_n(s)$$

$Eval(s)$ = material + mobility + king-safety + center-control

material = $10^{100}(K - K') + 9(Q - Q') + 5(R - R') + 3(B - B' + N - N') + 1(P - P')$

mobility = $0.1(\text{num-legal-moves} - \text{num-legal-moves}')$


Example for Evaluation Function Design




■ A danger of re-planning agents!

- He knows his score will go up by eating the dot now (west, east)
- He knows his score will go up just as much by eating the dot later (east, west)
- There are no point-scoring opportunities after eating the dot (within the horizon, two here)
- Therefore, waiting seems just as good as eating: he may go east, then back west in the next round of replanning!


Game Tree Pruning



Minimax Example




Minimax Pruning




Alpha-Beta Pruning

- General configuration (MIN version)
 - ➔ ▪ We're computing the MIN-VALUE at some node n
 - ➔ ▪ We're looping over n 's children
 - ➔ ▪ n 's estimate of the childrens' min is decreasing
 - ➔ ▪ Who cares about n 's value? MAX
 - ➔ ▪ Let a be the best value that MAX can get at any choice point along the current path from the root
 - ➔ ▪ If n becomes worse than a , MAX will avoid it, so we can stop considering n 's other children (it's already bad enough that it won't be played)
- MAX version is symmetric



Alpha-Beta Pruning

- General configuration (MAX version)
 - ➔ ▪ We're computing the MIN-VALUE at some node n
 - ➔ ▪ We're looping over n 's children
 - ➔ ▪ n 's estimate of the childrens' min is decreasing
 - ➔ ▪ Who cares about n 's value? MAX
 - ➔ ▪ Let a be the best value that MAX can get at any choice point along the current path from the root
 - ➔ ▪ If n becomes worse than a , MAX will avoid it, so we can stop considering n 's other children (it's already bad enough that it won't be played)
- MAX version is symmetric



Alpha-Beta Implementation


α : MAX's best option on path to root
 β : MIN's best option on path to root

```
def max-value(state,  $\alpha$ ,  $\beta$ ):  
    initialize  $v = -\infty$   
    for each successor of state:  
         $v = \max(v, \text{value}(\text{successor}, \alpha, \beta))$   
        if  $v \geq \beta$  return  $v$   
         $\alpha = \max(\alpha, v)$   
    return  $v$ 
```

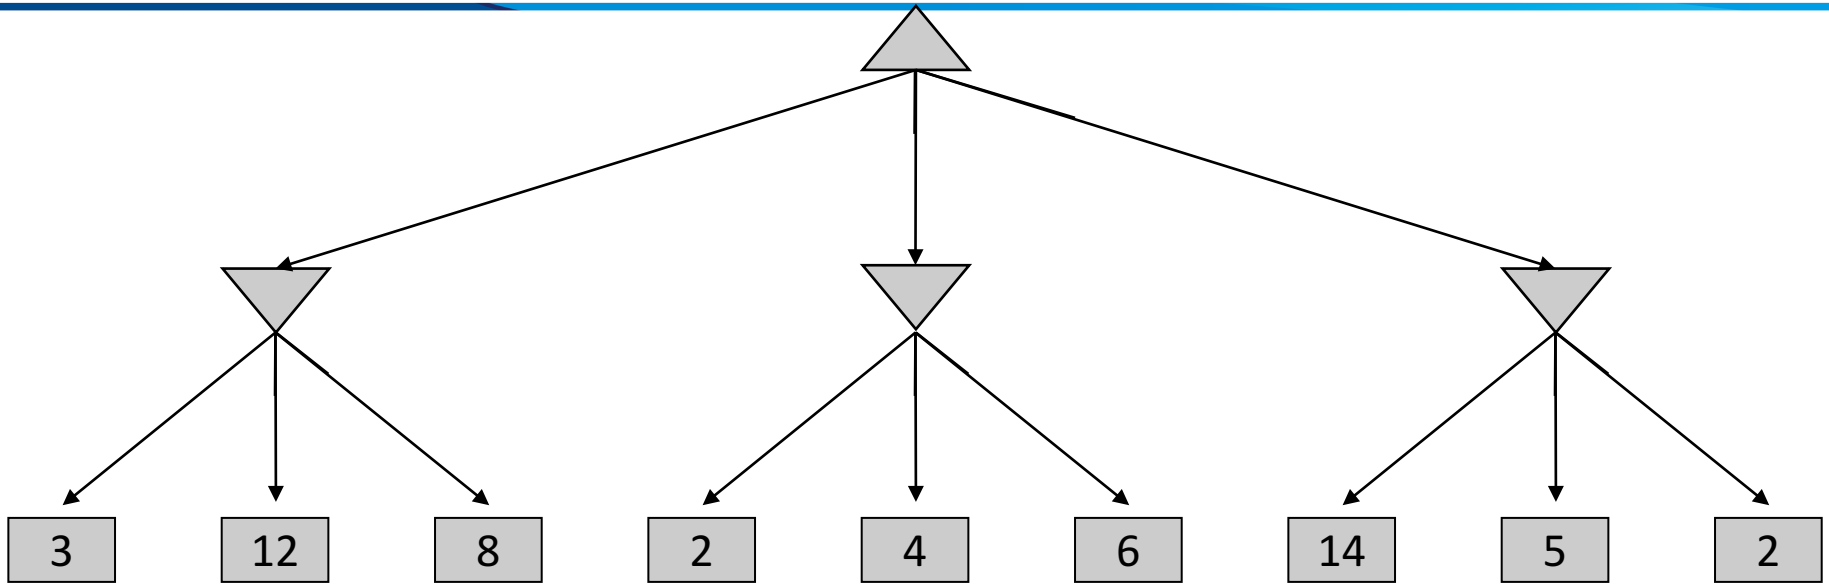
```
def min-value(state,  $\alpha$ ,  $\beta$ ):  
    initialize  $v = +\infty$   
    for each successor of state:  
         $v = \min(v, \text{value}(\text{successor}, \alpha, \beta))$   
        if  $v \leq \alpha$  return  $v$   
         $\beta = \min(\beta, v)$   
    return  $v$ 
```

Alpha-Beta Pruning Properties

- This pruning has **no effect** on minimax value computed for the root!
- Minimax values of intermediate nodes might be wrong
 - Important: children of the root may have the wrong value
 - So the most naïve version won't let you do action selection
- Good child ordering improves effectiveness of pruning
- With “perfect ordering”:
 - Time complexity drops to $O(b^{m/2})$
 - Doubles solvable depth!
 - Full search of, e.g. chess, is still hopeless...



Minimax Example



```
def max-value(state,  $\alpha$ ,  $\beta$ ):  
    initialize  $v = -\infty$   
    for each successor of state:  
         $v = \max(v, \text{value}(\text{successor}, \alpha, \beta))$   
        if  $v \geq \beta$  return  $v$   
         $\alpha = \max(\alpha, v)$   
    return  $v$ 
```

```
def min-value(state,  $\alpha$ ,  $\beta$ ):  
    initialize  $v = +\infty$   
    for each successor of state:  
         $v = \min(v, \text{value}(\text{successor}, \alpha, \beta))$   
        if  $v \leq \alpha$  return  $v$   
         $\beta = \min(\beta, v)$   
    return  $v$ 
```

Alpha-Beta Quiz 2

