


复旦大学大数据学院
School of Data Science, Fudan University

魏忠钰

Particle Filters and Applications of HMMs

May 16th, 2018


Particle Filtering



Particle Filtering


- Filtering: approximate solution
- Sometimes $|X|$ is too big to use exact inference
 - $|X|$ may be too big to even store $B(X)$
 - E.g. X is continuous
- Solution: approximate inference
 - Track samples of X , not all values
 - Samples are called particles
 - Time per step is linear in the number of samples
 - But: number needed may be large
 - In memory: list of particles, not states
- This is how robot localization works in practice
- Particle is just new name for sample

0.0	0.1	0.0
0.0	0.0	0.2
0.0	0.2	0.5



Representation: Particles

- Our representation of $P(X)$ is now a list of N particles (samples)
 - Generally, $N \ll |X|$
 - Storing map from X to counts would defeat the point
- $P(x)$ approximated by number of particles with value x
 - So, many x may have $P(x) = 0$!
 - More particles, more accuracy
- For now, all particles have a weight of 1



Particles:

(3,3)

(2,3)

(3,3)

(3,2)

(3,3)

(3,2)

(1,2)

(3,3)

(3,3)

(2,3)

Particle Filtering: Elapse Time


- Each particle is moved by sampling its next position from the transition model

$$x' = \text{sample}(P(X'|x))$$

- This is like prior sampling – samples' frequencies reflect the transition probabilities
 - Here, most samples move clockwise, but some move in another direction or stay in place
- This captures the passage of time
 - If enough samples, close to exact values before and after (consistent)


Particles:

(3,3)
(2,3)
(3,3)
(3,2)
(3,3)
(3,2)
(1,2)
(3,3)
(3,3)
(2,3)



Particles:

(3,2)
(2,3)
(3,2)
(3,1)
(3,3)
(3,2)
(1,3)
(2,3)
(3,2)
(2,2)



Particle Filtering: Observe

- Similar to likelihood weighting, re-weight samples based on the evidence


$$w(x) = P(e|x)$$

$$B(X) \propto P(e|X)B'(X)$$

- As before, the probabilities don't sum to one, since all have been re-weighted (in fact they now sum to (N times) an approximation of $P(e)$)


Particles:

(3,2)
(2,3)
(3,2)
(3,1)
(3,3)
(3,2)
(1,3)
(2,3)
(3,2)
(2,2)



Particles:

(3,2) w=.9
(2,3) w=.2
(3,2) w=.9
(3,1) w=.4
(3,3) w=.4
(3,2) w=.9
(1,3) w=.1
(2,3) w=.2
(3,2) w=.9
(2,2) w=.4




Particle Filtering: Resample

- Rather than tracking weighted samples, we resample
- N times, we choose from our weighted sample distribution (i.e. draw with replacement)
- This is equivalent to renormalizing the distribution
- Now the update is complete for this time step, continue with the next one


Particles:

(3,2) $w=.9$
(2,3) $w=.2$
(3,2) $w=.9$
(3,1) $w=.4$
(3,3) $w=.4$
(3,2) $w=.9$
(1,3) $w=.1$
(2,3) $w=.2$
(3,2) $w=.9$
(2,2) $w=.4$




(New) Particles:

(3,2)
(2,2)
(3,2)
(2,3)
(3,3)
(3,2)
(1,3)
(2,3)
(3,2)
(3,2)





Recap: Particle Filtering

- Particles: track samples of states rather than an explicit distribution



Robot Localization


- In robot localization:
 - We know the map, but not the robot's position
 - Observations may be vectors of range finder readings
 - State space and readings are typically continuous (works basically like a very fine grid) and so we cannot store $B(X)$
 - Particle filtering is a main technique





Robot Mapping


- SLAM: Simultaneous Localization And Mapping
 - We do not know the map or our location
 - State consists of position AND map!
 - Main techniques: Kalman filtering (Gaussian HMMs) and particle methods



- **Smoothing**

- Computing the posterior distribution over a *past* state, given all evidence up to the present.
- $P(X_k | e_{1:t})$

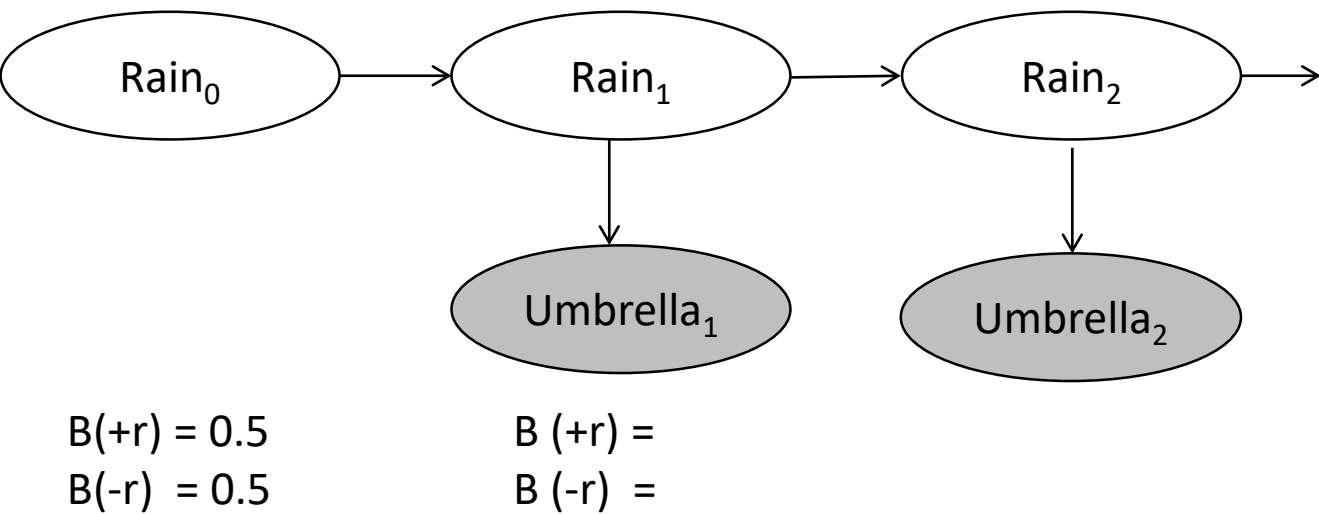
Smoothing



$$\begin{aligned} P(X_k | \mathbf{e}_{1:t}) &= P(X_k | \mathbf{e}_{1:k}, \mathbf{e}_{k+1:t}) \\ &\propto P(X_k | \mathbf{e}_{1:k}) P(\mathbf{e}_{k+1:t} | X_k, \mathbf{e}_{1:k}) \\ &= P(X_k | \mathbf{e}_{1:k}) P(\mathbf{e}_{k+1:t} | X_k) \end{aligned}$$

$$\begin{aligned} P(\mathbf{e}_{k+1:t} | X_k) &= \sum_{\mathbf{x}_{k+1}} P(\mathbf{e}_{k+1:t} | X_k, \mathbf{x}_{k+1}) P(\mathbf{x}_{k+1} | X_k) \\ &= \sum_{\mathbf{x}_{k+1}} P(\mathbf{e}_{k+1:t} | \mathbf{x}_{k+1}) P(\mathbf{x}_{k+1} | X_k) \\ &= \sum_{\mathbf{x}_{k+1}} P(\mathbf{e}_{k+1}, \mathbf{e}_{k+2:t} | \mathbf{x}_{k+1}) P(\mathbf{x}_{k+1} | X_k) \\ &= \sum_{\mathbf{x}_{k+1}} P(\mathbf{e}_{k+1} | \mathbf{x}_{k+1}) P(\mathbf{e}_{k+2:t} | \mathbf{x}_{k+1}) P(\mathbf{x}_{k+1} | X_k) \end{aligned}$$

Example: Weather HMM



R_t	R_{t+1}	$P(R_{t+1} R_t)$
+r	+r	0.7
+r	-r	0.3
-r	+r	0.3
-r	-r	0.7

R_t	U_t	$P(U_t R_t)$
+r	+u	0.9
+r	-u	0.1
-r	+u	0.2
-r	-u	0.8

Example: Weather HMM


Filtering

$$B(+r) = 0.5$$

$$B(-r) = 0.5$$

$$B(+r) = 0.818$$

$$B(-r) = 0.182$$



R_t	R_{t+1}	$P(R_{t+1} R_t)$
+r	+r	0.7
+r	-r	0.3
-r	+r	0.3
-r	-r	0.7

R_t	U_t	$P(U_t R_t)$
+r	+u	0.9
+r	-u	0.1
-r	+u	0.2
-r	-u	0.8

Smoothing

$$B(+r) = 0.5$$

$$B(-r) = 0.5$$

$$f: B(+r) = 0.818$$

$$f: B(-r) = 0.182$$

$$b: B(+r) = 0.69$$

$$b: B(-r) = 0.41$$



$$B(+r) = 0.883$$

$$B(-r) = 0.117$$

- **Most Likely Explanation**

- Given a sequence of observations, find the sequence of states that is most likely to have generated those observations.


$$\arg \max_{x_{1:t}} P(x_{1:t} | e_{1:t})$$

Most Likely Explanation




HMMs: MLE Queries

- HMMs defined by
 - States X
 - Observations E
 - Initial distribution: $P(X_1)$
 - Transitions: $P(X|X_{-1})$
 - Emissions: $P(E|X)$
- New query: most likely explanation: $\arg \max_{x_{1:t}} P(x_{1:t}|e_{1:t})$
- New method: the Viterbi algorithm




Viterbi

- Graph of states and transitions over time



- Each arc represents some transition $x_{t-1} \rightarrow x_t$
- Each arc has weight $P(x_t|x_{t-1})P(e_t|x_t)$
- Each path is a sequence of states
- The product of weights on a path is that sequence's probability along with the evidence
- Forward algorithm computes sums of paths, Viterbi computes best paths

Forward / Viterbi Algorithms



Forward Algorithm (Sum)

$$f_t[x_t] = P(x_t, e_{1:t})$$

$$= P(e_t|x_t) \sum_{x_{t-1}} P(x_t|x_{t-1}) f_{t-1}[x_{t-1}]$$

Viterbi Algorithm (Max)

$$m_t[x_t] = \max_{x_{1:t-1}} P(x_{1:t-1}, x_t, e_{1:t})$$

$$= P(e_t|x_t) \max_{x_{t-1}} P(x_t|x_{t-1}) m_{t-1}[x_{t-1}]$$

▪ Filtering

- Computing the **belief state**—the posterior distribution over the most recent state—given all evidence to date.

- $P(X_t | e_{1:t})$

▪ Prediction

- Computing the posterior distribution over the *future* state, given all evidence to date.

- $P(X_{t+k} | e_{1:t})$

▪ Smoothing

- Computing the posterior distribution over a *past* state, given all evidence up to the present.

- $P(X_k | e_{1:t})$

▪ Most Likely Explanation

- Given a sequence of observations, find the sequence of states that is most likely to have generated those observations.

$$\arg \max_{x_{1:t}} P(x_{1:t} | e_{1:t})$$